

CSE 310

Data Structures and Algorithms

Instructor: Hanghang Tong
Assistant Professor
CIDSE

Course materials available through Blackboard on MyASU



Lecture Overview

- Discussion on syllabus and logistics
 - Prerequisites, topics to cover, assessment
- Introduction to Algorithms & Data Structures
 - What is an algorithm; What do we care about an algorithm
 - Examples of algorithms

About This Course

- This is a core course of our CSE programs.
 - It is typically viewed as one of the most difficult courses by CSE students here.
 - Your being a straight-A student *thus far* does not mean you can have an easy pass, let alone an easy A in this course.
 - But it is also one of those few courses that makes you distinctively “computer science”-ish (versus other engineering majors)
 - Try to do this yourself afterwards: search on the Internet about “Google interview questions”

Staff

- Instructor: Hanghang Tong
 - Doing research on data mining, big data analytics → all involving designing sophisticated algorithms
 - Office: BY416, Office Hours: M/W 10-11am or other times by appointment.
 - TAs: (All the office hours by TAs will be at CenterPoint 114)
 - Joydeep Banerjee (joydeep.banerjee@asu.edu), OH: Tues/Thurs 1-2pm
 - Jun Shen (jun.shen.1@asu.edu), OH: Tues 10-11am, Wedn 1-2pm
 - Huijun Wu (huijun.wu@asu.edu), OH: Mon 1:15-2:15pm, Wedn: 10:30-11:30am
 - Grader: Rongyu Lin (No OH)
-
- Between the TAs and me, we have a total of 8 hours of office hours per week for helping you outside the class hours.
 - Additional Resource for Help: Discussion Board on BB.

Prerequisites

- For CS/CSE students, grades of at least C in CSE 220 or CSE 240 and at least D in MAT 243. For CMS students, grades of at least C in CSE 210 and at least D in MAT 243 or MAT 300.

Programming experience in C or C++ is expected.

- Read the Appendix A of the textbook: the appendix should read like a review of known materials rather than new knowledge to you (except perhaps C, which is on probabilities, but you should at least be taking a probability course now).
- Try the self-review test in BB
 - *Otherwise, this course may not be right for you.*

- Proficiency in C/C++ programming is required

- Class project will require programming in C/C++, in Linux.
- This course will NOT teach you how to program in C/C++; You are assumed to know that already. *Otherwise, this course may not be right not for you.*

Course Information

- Course materials will be available only through the Blackboard on MyASU.
- Required textbooks:
 - ***Introduction to Algorithms***, Cormen, T. H., Leiserson, C.E., Rivest, R.L., and Stein, C., The MIT Press, Cambridge Massachusetts, 3rd edition, 2009.
 - This is a ***required*** book. You need to have it handy immediately. Cannot use “no book” as an excuse for late submission of assignments.
 - I’m not aware of the existence of a “desk copy” either. ➔ ***Purchase your copy now if you haven’t done so.***
 - It is a book worth keeping as a reference for your job interviews and even your future jobs.

Course Information

- Lecture notes will be posted before each class
 - May be updated throughout the semester. You need to update your version accordingly.
 - Lecture notes may **NOT** include examples that are worked out on the whiteboard.
 - Good attendance is key for efficient and effective study
- Lecture Note Takers:

A notetaker is needed for this class. In exchange for providing this service, the selected notetaker will be paid a stipend of \$25.00 per credit, i.e. total of \$75 for a 3 credit class OR if preferred, awarded a letter of community service at the end of the semester. If you take clear, concise notes and are willing to share a copy, please follow up with me after class to receive instructions on how to sign up online.

Course Information

- Homework assignments (including the class projects) and supplemental reading materials will all be posted on the Blackboard.
 - **Late submissions of the assignments will not be accepted.**
- Again, the class projects will require programming work. All programming work needs to be completed in C/C++, in Linux
 - Details will be announced soon on BB

Assessment

- Homework 10%: 4 of equal weight.
 - No late submission will be accepted.
- Projects: 15%, 3 equal weight, **all in c/c++ only on general.asu.edu (Linux)**
 - No late submission will be accepted.
- Midterm Exams 30% : there will be **two** of them, closed book.
 - Each exam covers only the period from the previous exam until the current exam.
 - The two exams amount to 30% total (equal weight).
- Final Exam 45%: 9:50-11:40am, December 7th, comprehensive, closed book
- **No make-up exam will be given unless it is for genuine, verifiable emergency.**

Assessment – Key Dates

Assignments	Out	Due	Projects	Out	Milestone	Due	Midterms	Scheduled
A1	08/24	09/11	P1	08/24	09/09	09/23	M1	09/30
A2	09/16	10/09	P2	09/23	10/16	10/28	M2	11/09
A3	10/14	11/06	P3	10/28	11/20	12/04		
A4	11/06	11/25						

Assessment

- The following cutoffs represent what will be *likely* used to generate the letter grades:

A+ $\geq 97\%$	A $\geq 90\%$ & $< 97\%$	A- $\geq 87\%$ & $< 90\%$
B+ $\geq 84\%$ & $< 87\%$	B $\geq 80\%$ & $< 84\%$	B- $\geq 77\%$ & $< 80\%$
C+ $\geq 74\%$ & $< 77\%$	C $\geq 70\%$ & $< 74\%$	C- $\geq 67\%$ & $< 70\%$
	D $\geq 60\%$ & $< 67\%$	E $< 60\%$

- The above cutoffs are tentatively and may be adjusted *slightly*; However, there will be *no general curve-fitting* in assigning the final grades.

Major Topics & Tentative Schedule

- | | |
|----|---|
| 1 | Syllabus & Intro, Asymptotic Notation, |
| 2 | Merge Sort, Quick Sort |
| 3 | Heaps, Priority Queues |
| 4 | Sorting Lower Bound, Linear sort algorithms (briefly);
Selection |
| 5 | Selection; Hash Tables |
| 6 | Hash Tables; Binary search tree |
| 7 | Red-Black Trees |
| 8 | Disjoint Sets |
| 9 | Matrix Multiplication |
| 10 | Longest common sequence; Depth First Search |
| 11 | Breadth First Search; Topological Sort & SCC |
| 12 | Minimum Spanning Trees; Shortest Paths |

Roughly, each line corresponds to 1 week.

Questions or Comments

Introduction to Algorithms


- What is an algorithm: An algorithm is a well-defined computational procedure that takes some input, produces some output, and then terminate.
- An example of an algorithm: design a computational procedure for sorting a given set of n number. E.g., $n = 5$ for the example below.

Given input: 5, 2, 4, 10, 7

Desired output: 2, 4, 5, 7, 10

Your fancy (computational) way of generating the desired output for any given input → An algorithm

Analogy to sorting cards ...



- Think about what you would do
- There may be many ways, but one intuitive *procedure* is to always *insert* a card into the right place
 - Works naturally when you draw a stack of cards one by one and put them into your hand.
 - Similarly, if you sort all cards in your hand → “*in place*” sorting

Computationally represent the procedure

- Pseudocode: Liberal use of English; Use of indentation for block structure; Omission of error handling (e.g., check if an array is empty).

```
// Pseudocode for Insertion Sort
```

```
INSERTION-SORT(A)
```

```
for  $j = 2$  to  $A.length$  //  $A.length = n$ 
```

```
    key =  $A[j]$ 
```

```
    // Insert  $A[j]$  into sorted sequence  $A[1 \dots j-1]$ 
```

```
     $i = j - 1$ 
```

```
    while  $i > 0$  and  $A[i] > key$ 
```

```
         $A[i+1] = A[i]$ 
```


```
         $i = i - 1$ 
```

```
     $A[i+1] = key$ 
```



So, how good is it?

- Does it work? Is it fast enough? Does it require a lot of memory?
→ Analysis of algorithms.



We want to answer these questions **independent** of machines, programming languages, etc.

→ **RAM (Random Access Machine)**: Model of computation

- Basic instruction can be performed in **one unit of time (time step)**.

Complexity of An Algorithm

- Time complexity: **Running time of an algorithm** in terms of the **size of the input** to the algorithm.
- Space complexity: measured using the number of memory units required for executing the procedure.
- Both time complexity and space complexity are in general a function of the input size, typically denoted n .
 - **Sorting 100 elements takes more time, as well as requires more memory, than sorting 3 elements**

Let's analyze our Insertion Sort algorithm

-- First look at how it works on a specific input.

Example:

3 9 6 1 2 9 should be inserted after 3 – no change

3 9 6 1 2 6 should be inserted between 3 and 9

3 6 9 1 2 1 should be inserted before 3

1 3 6 9 2 2 should be inserted between 1 and 3

1 2 3 6 9

<u>times</u>	<u>cost</u>
n	c_1
$n-1$	c_2
0	$c_3=0$
$n-1$	c_4
$\sum_{j=2}^n t_j$	c_5
$\sum_{j=2}^n (t_j-1)$	c_6
$\sum_{j=2}^n (t_j-1)$	c_7
$n-1$	c_8

// Pseudocode for Insertion Sort

INSERTION-SORT(A)

for $j = 2$ to A.length // A.length = n

 key = A[j]

 // Insert A[j] into sorted sequence A[1 ... j-1]

$i = j - 1$

while $i > 0$ and A[i] > key

 A[i+1] = A[i]

$i = i - 1$

 A[i+1] = key

t_j = the number of times the while loop test is executed for that value of j .

$t_j - 1$ = # of elements in A[1...j-1] that are $> A[j]$

For the insertion sort in the previous pages, its running time is:

$$T(n) = c_1n + c_2(n-1) + c_4(n-1) + c_5 \sum_{j=2}^n t_j + c_6 \sum_{j=2}^n (t_j - 1) \\ + c_7 \sum_{j=2}^n (t_j - 1) + c_8(n-1)$$

$T(n)$ depends not just on n , but also the kinds of input arrays:

BEST-CASE: array is sorted.

In this case, $t_j = 1$ for all j

The running time becomes:

$$c_1n + c_2(n-1) + c_4(n-1) + c_5(n-1) + c_8(n-1) \\ = (c_1 + c_2 + c_4 + c_5 + c_8)n - (c_2 + c_4 + c_5 + c_8) \\ \rightarrow \text{A linear function of } n.$$

WORST-CASE: array is in reverse order

In this case, $t_j = j$ for all j

The running time becomes:

$$\begin{aligned} & c_1 n + c_2(n-1) + c_4(n-1) + c_5 \sum_{j=2}^n j + c_6 \sum_{j=2}^n (j-1) + c_7 \sum_{j=2}^n (j-1) + c_8(n-1) \\ &= c_1 n + c_2(n-1) + c_4(n-1) + c_5 [n(n+1)/2 - 1] + c_6 [n(n-1)/2] \\ &\quad + c_7 [n(n-1)/2] + c_8(n-1) \\ &= (c_5/2 + c_6/2 + c_7/2) n^2 + (c_1 + c_2 + c_4 + c_5/2 - c_6/2 - c_7/2 + c_8) n - (c_2 + c_4 + c_5 + c_8) \end{aligned}$$

➔ A quadratic function of n .

Note: $\sum_{j=2}^n j = n(n+1)/2 - 1$ and $\sum_{j=2}^n (j-1) = n(n-1)/2$

--- See Appendix A of the textbook for more summation.

- What's the big deal about the difference between “a linear function” and “a quadratic function” of n ?
 - $n = 10 \rightarrow n^2 = 100$, $n = 1000 \rightarrow n^2 = 1,000,000$, ...
- There are worse algorithms whose complexity may be of n^3 or even 2^n
 - The difference will be even more! \rightarrow Eventually render some algorithms with a high complexity practically useless except for very small n .

Data Structure

- The course title contains the term “Data Structure”: why is it there?
- A data structure is a way to store and organize data in order to facilitate access and modifications.
 - Think about how you would compile a dictionary (to facilitate looking up any entry)
 - Think about how you would organize all songs on your PC (to facilitate retrieving a song of your choice)
- Efficient algorithms are often designed together with accompanying data structures
 - Graph-based algorithms

Summary & Reading Assignment

- What is an algorithm?
- Some basic ideas in analyzing an algorithm
- An examples: insertion sort
 - ➔ Read to review: Chapter 1, Chapter 2.1, 2.2
- Next time: formal analysis techniques: asymptotic notations
 - ➔ Read to prepare: Chapter 3